

HCMC University Of Technology
Faculty of Computer Science Engineering



course: Operating Systems

Assignment Simple Operating System

Instructor: Nguyễn Lê Duy Lai - Nguyễn Mạnh Thìn

Class: CC07

Supervised by:

Lương Thị Minh Oanh - 1950031

Nguyễn Đặng Tứ Hải - 1952045

Date submitted: 6th May 2022

Contents

1 Scheduler

- 1.1 Question
- 1.1.1 Question
- 1.1.2 Answer
- 1.2 code
- 1.3 Results

2 Memory Management

- 2.1 Question
- 2.1.1 Question
- 2.1.2 Answer
- 2.2 Implementation
- 2.2.1 Search for page table from the a segment table
- 2.2.2 Translate virtual address into physical address
- 2.2.3 Memory allocation
- 2.2.4 Free memory
- 2.3 Test run

3 Put all together

1 Scheduler

1.1 Question

1.1.1 Question

What is the advantage of using priority feedback queue in comparison with other scheduling algorithms you have learned?

1.1.2 Answer

The Priority Feedback Queue (PFQ) algorithm uses the ideas of a number of other algorithms including Priority Scheduling - each process has a priority to execute, Multilevel Queue algorithm - uses multi-level process queues, Round Robin algorithm - uses quantum time for processes to execute. Here are other scheduling algorithms we have learned:

First Come First Served (FCFS)

Shortest Job First (SJF)

Priority Scheduling (PS)

Multilevel Queue Scheduling (MLQS)

Round Robin (RR)

Multilevel Feedback Queue (MLFQ)

Shortest Remaining Time First (SRTF)

Priority Feedback Queue (PFQ) use two queue: `ready_queue` and `run_queue`:

- `ready_queue`: This queue contains processes that will go to the CPU. And this queue's priority is higher than `run_queue`.
- `run_queue`: This queue contains the waiting processes to be executed after their time slot runs off. The processes in this queue only move to the `ready_queue` when the `ready_queue` is free.
- These two queues both have priority based on the process priority.

Advantages of using priority feedback queue:

- Avoiding starvation: Using the idea of RR with a quantum time. Each process is served by the CPU for a fixed time quantum. After that, the process will be moved to `run_queue` and the lower-priority processes can use the CPU. Therefore, we can avoid the problem that some processes have to wait for too long (starvation) in some methods such as:

+ SJF (Shortest Job First): Longer processes will have more waiting time, eventually they'll suffer starvation.

+ SRTF (Shortest Remaining Time First), Priority based scheduling: Processes with low priority may not be allowed to use the CPU.

- Using two queues, like the idea of MLQS and MLFQ, where two queues are cycle through processes until the process is completed, increasing the response time for processes (the Processes with lower priority that come later can still be executed before processes with higher priority after you have finished your slot).

- Processes are executed based on priority: Important process can be executed first by assigning its priority higher than those processes waiting in `ready_queue`. It helps us to control which process will be executed first, and which will be executed later. Some scheduling algorithms such as: FCFS (First Come First Served), SJF (Shortest Job First), SRTF (Shortest Remaining Time First), RR (Round Robin) can not do this

1.2 code

Firstly, `enqueue()` in `queue.c`: we only need to put a new process to the end of the queue `[q]` if it is not full, otherwise, we return and do nothing. Here is the implementation:

```
void enqueue(struct queue_t * q, struct pcb_t * proc) {
    /* TODO: put a new process to queue [q] */
    if (q->size < MAX_QUEUE_SIZE) {
        q->proc[q->size] = proc;
        q->size++;
    }
}
```

Secondly, `dequeue()` in `queue.c`: we look for the element of the queue `[q]` with the highest priority. If the queue `[q]` is empty, we return NULL. Otherwise, we assign its position into variable `position`. Then we shift every element from `position + 1` to the end of the queue `[q]` to its front one position. Finally, we get the process with the

highest priority is the variable result equal the element at the position in the queue [q]. Here is the implementation:

```
struct pcb_t * dequeue(struct queue_t * q) {
    /* TODO: return a pcb whose priority is the highest
     * in the queue [q] and remember to remove it from q
     * */
    if (q->size == 0) return NULL;
    struct pcb_t* temp_proc = NULL;
    /* For finding the highest priority process in queue */
    int current_max_priority = q->proc[0]->priority;
    temp_proc = q->proc[0];
    /* For saving index while searching highest priority process */
    int found_index = 0;
    int i;
    for (i = 1; i < q->size; i++) {
        if (q->proc[i]->priority > current_max_priority) {
            temp_proc = q->proc[i];
            current_max_priority = q->proc[i]->priority;
            found_index = i;
        }
    }
    /* Re-Update queue */
    for (i = found_index; i < q->size - 1; i++)
        q->proc[i] = q->proc[i + 1];
    q->proc[q->size - 1] = NULL;
    --q->size;
    return temp_proc;
}
```

Finally, `get_proc` in `shed.c`: The task of the scheduler is to manage the updating of the processes that will be executed to the CPU. In this assignment, we just need to implement the function to find a process for the CPU to execute. In more details, with the function `get_proc()`, if the `ready_queue` is empty, we move all processes waiting in `run_queue` back to `ready_queue`. After that, we use `dequeue()` function to return the process that has the highest priority from `ready_queue`.

Here is the implementation:

```
struct pcb_t * get_proc(void) {
    struct pcb_t* proc = NULL;
```

```
/*TODO: get a process from [ready_queue]. If ready queue
* is empty, push all processes in [run_queue] back to
* [ready_queue] and return the highest priority one.
* Remember to use lock to protect the queue.
* */
pthread_mutex_lock(&queue_lock);
/* If ready queue is empty,
* move all processes in run queue to ready queue
*/
if (ready_queue.size == 0) {
    int i;
    for (i = 0; i < MAX_QUEUE_SIZE; i++){
        ready_queue.proc[i] = run_queue.proc[i];
        run_queue.proc[i] = NULL;
    }
    ready_queue.size = run_queue.size;
    run_queue.size = 0;
}
proc = dequeue(&ready_queue);
pthread_mutex_unlock(&queue_lock);
return proc;
}
```

1.3 Results

To open Terminal in Makefile, run "make sched" to check code, and run "make test_sched" to run program. The result of sched_0 (an example of different results):

```
Time slot    0
Loaded a process at input/proc/s0, PID: 1
CPU 0: Dispatched process 1
Time slot    1
Time slot    2
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 1
Time slot    3
```


```
Time slot 4
Loaded a process at input/proc/s1, PID: 2
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 2
Time slot 5
Time slot 6
CPU 0: Put process 2 to run queue
CPU 0: Dispatched process 2
Time slot 7
Time slot 8
CPU 0: Put process 2 to run queue
CPU 0: Dispatched process 1
Time slot 9
Time slot 10
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 2
Time slot 11
Time slot 12
CPU 0: Put process 2 to run queue
CPU 0: Dispatched process 1
Time slot 13
Time slot 14
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 2
Time slot 15
CPU 0: Processed 2 has finished
CPU 0: Dispatched process 1
Time slot 16
Time slot 17
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 1
Time slot 18
Time slot 19
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 1
Time slot 20
Time slot 21
CPU 0: Put process 1 to run queue
```

1.3 Results

/

```
CPU 0: Dispatched process 1
Time slot 22
CPU 0: Processed 1 has finished
CPU 0 stopped
```

We use Gantt chart to illustrate this test ./os sched_0:



In this test, CPU runs on 2 processes which are P1 and P2 in 23 time slot as indicated on the Gantt chart above.


And the result of sched_1:

```
Time slot 0
Loaded a process at input/proc/s0, PID: 1
Time slot 1
CPU 0: Dispatched process 1
Time slot 2
Time slot 3
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 1
Time slot 4
Loaded a process at input/proc/s1, PID: 2
Time slot 5
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 2
Time slot 6
Loaded a process at input/proc/s2, PID: 3
Time slot 7
CPU 0: Put process 2 to run queue
CPU 0: Dispatched process 3
Loaded a process at input/proc/s3, PID: 4
Time slot 8
Time slot 9
CPU 0: Put process 3 to run queue
CPU 0: Dispatched process 4
Time slot 10
```


Time slot 11
CPU 0: Put process 4 to run queue
CPU 0: Dispatched process 2
Time slot 12
Time slot 13
CPU 0: Put process 2 to run queue
CPU 0: Dispatched process 3
Time slot 14
Time slot 15
CPU 0: Put process 3 to run queue
CPU 0: Dispatched process 1
Time slot 16
Time slot 17
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 4
Time slot 18
Time slot 19
CPU 0: Put process 4 to run queue
CPU 0: Dispatched process 2
Time slot 20
Time slot 21
CPU 0: Put process 2 to run queue
CPU 0: Dispatched process 3
Time slot 22
Time slot 23
CPU 0: Put process 3 to run queue
CPU 0: Dispatched process 1
Time slot 24
Time slot 25
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 4
Time slot 26
Time slot 27
CPU 0: Put process 4 to run queue
CPU 0: Dispatched process 2
Time slot 28
CPU 0: Processed 2 has finished
CPU 0: Dispatched process 3

Time slot 29
Time slot 30
CPU 0: Put process 3 to run queue
CPU 0: Dispatched process 1
Time slot 31
Time slot 32
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 4
Time slot 33
Time slot 34
CPU 0: Put process 4 to run queue
CPU 0: Dispatched process 3
Time slot 35
Time slot 36
CPU 0: Put process 3 to run queue
CPU 0: Dispatched process 1
Time slot 37
Time slot 38
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 4
Time slot 39
Time slot 40
CPU 0: Put process 4 to run queue
CPU 0: Dispatched process 3
Time slot 41
Time slot 42
CPU 0: Processed 3 has finished
CPU 0: Dispatched process 1
Time slot 43
Time slot 44
CPU 0: Put process 1 to run queue
CPU 0: Dispatched process 4
Time slot 45
CPU 0: Processed 4 has finished
CPU 0: Dispatched process 1
Time slot 46
CPU 0: Processed 1 has finished
CPU 0 stopped

We also use Gantt chart to illustrate this test ./os sched_1:



In this test, CPU runs on 4 processes which are P1, P2, P3 and P4 in 46 time slot as indicated on Gantt chart above.

2 Memory Management

2.1 Question

2.1.1 Question

What is the advantage and disadvantage of segmentation with paging.

2.1.2 Answer

Advantages:

- Saving memory
- Page size is limited by segment size
- Simplify memory allocation
- Overcome external fragmentation

Disadvantages:

- Causing internal fragmentation
- Complex level is higher than segmented paging
- Page table needs to store continuously

2.2 Implementation

2.2.1 Search for page table from the a segment table

In this assignment, address is represented by 20 bits, we use the first 5 bits for segment index, the next 5 bits for page index and the last 10 bits for offset. We

perform `get_page_table()` function in `mem.c` as below:

```
/* Search for page table table from the a segment table */
static struct page_table_t* get_page_table(

    addr_t index, //Segment level index
    struct seg_table_t* seg_table) { //First level table

    /*
    * TODO: Given the Segment index [index], you must go through each
    * row of the segment table [seg_table] and check if the v_index
    * field of the row is equal to the index
    *
    * */
    if (seg_table == NULL) return NULL;
    int i;
    for (i = 0; i < seg_table->size; i++) {
        // Enter your code here
        if (seg_table->table[i].v_index == index) {
            return seg_table->table[i].pages;
        }
    }
    return NULL;
}
```

From first 5 bits, which is index of segment table, if we find table that has `v_index` equal to index of transformed table, we will return corresponding page .

2.2.2 Translate virtual address into physical address

Each address contains 20 bits with mentioned organization. To create physical address, we take first 10 bits (segment and page) and connect it with last 10 bits (offset). Each `page_table_t` store elements with `p_index` of first 10 bits. We need to shift left 10 bits or multiple with `PAGE_SIZE` then add with offset. `translate()` function is implemented as below:

```
/* Translate virtual address to physical address. If [virtual_addr] is valid,
* return 1 and write its physical counterpart to [physical_addr].
* Otherwise, return 0 */
static int translate(addr_t virtual_addr, // Given virtual address
```

```

addr_t* physical_addr, // Physical address to be returned
struct pcb_t* proc) { // Process uses given virtual address

    /* Offset of the virtual address */
    addr_t offset = get_offset(virtual_addr);
    /* The first layer index */
    addr_t first_lv = get_first_lv(virtual_addr);
    /* The second layer index */
    addr_t second_lv = get_second_lv(virtual_addr);

    /* Search in the first level */
    struct page_table_t* page_table = NULL;
    page_table = get_page_table(first_lv, proc->seg_table);
    if (page_table == NULL) {
        return 0;
    }

    int i;
    for (i = 0; i < page_table->size; i++) {
        if (page_table->table[i].v_index == second_lv) {
            /* TODO: Concatenate the offset of the virtual address
            * to [p_index] field of page_table->table[i] to
            * produce the correct physical address and save it to
            * [*physical_addr] */
            *physical_addr = page_table->table[i].p_index * PAGE_SIZE + offset;
            return 1;
        }
    }
    return 0;
}

```

2.2.3 Memory allocation

Firstly, we check if both virtual and physical address are ready to use. In physical address, we find and count empty pages, if the quantity of empty pages exceeds the limitation, physical address is ready to use. In virtual address, we use Break Point of process, if it satisfy RAM-size condition, virtual address is ready to use. We create

```

check_available_memory(uint32_t num_page, struct * pcb_t proc):

int check_available_memory(uint32_t num_pages, struct pcb_t *proc) {
    int num_of_valid_pages = 0;
    int i;
    for (i = 0; i < NUM_PAGES; i++) {
        if (_mem_stat[i].proc == 0) { // No process here
            num_of_valid_pages++;
            // If find enough space for current process size
            if (num_of_valid_pages == num_pages && proc->bp + num_pages * PAGE_
                SIZE <= RAM_SIZE) {
                return 1;
                break;
            }
        }
    }
    return 0;
}

```

if the result is 1, memory zone is enable, then we compile `alloc_mem()`

```

addr_t alloc_mem(uint32_t size, struct pcb_t* proc) {
    pthread_mutex_lock(&mem_lock);
    addr_t ret_mem = 0;
    /* TODO: Allocate [size] byte in the memory for the
     * process [proc] and save the address of the first
     * byte in the allocated memory region to [ret_mem].
     */

    uint32_t num_pages = (size % PAGE_SIZE == 0) ? size / PAGE_SIZE :
        size / PAGE_SIZE + 1; // Number of pages we will use

    int mem_avail = 0; // We could allocate new memory region or not?

    /* First we must check if the amount of free memory in
     * virtual address space and physical address space is
     * large enough to represent the amount of required
     * memory. If so, set 1 to [mem_avail].
     * Hint: check [proc] bit in each page of _mem_stat

```

```

    * to know whether this page has been used by a process.
    * For virtual memory space, check bp (break pointer).
    * */
mem_avail = check_available_memory(num_pages, proc);
if (mem_avail) {
    /* We could allocate new memory region to the process */
    ret_mem = proc->bp;
    proc->bp += num_pages * PAGE_SIZE;
    /* Update status of physical pages which will be allocated
       * to [proc] in _mem_stat. Tasks to do:
       * - Update [proc], [index], and [next] field
       * - Add entries to segment table page tables of [proc]
       * to ensure accesses to allocated memory slot is
       * valid. */
    int count_alloc_pages = 0;
    int prev_index = -1;
    addr_t virtual_addr;
    int first_lv, second_lv;
    int i;
    for (i = 0; i < NUM_PAGES; i++) {
        if (_mem_stat[i].proc) continue;

        // If found a free hole
        _mem_stat[i].proc = proc->pid;
        _mem_stat[i].index = count_alloc_pages;

        if (prev_index > -1)
            _mem_stat[prev_index].next = i;
        prev_index = i;

        int isFound = 0;
        struct seg_table_t *seg_table = proc->seg_table;
        // Seg_table is empty
        if (seg_table->table[0].pages == NULL)
            seg_table->size = 0;

        virtual_addr = ret_mem + (count_alloc_pages * PAGE_SIZE);
    }
}

```

```

        // Construct page tables
        first_lv = get_first_lv(virtual_addr);
        second_lv = get_second_lv(virtual_addr);
        int k;
        for (k = 0; k < seg_table->size; k++) {
            // If found a page table
            if (seg_table->table[k].v_index == first_lv) {
                // Allocate new slot in current page table
                struct page_table_t* cur_page_table = seg_table->table[k].pages;
                cur_page_table->table[cur_page_table->size].v_index = second_lv;
                cur_page_table->table[cur_page_table->size].p_index = i;
                isFound = 1;
                break;
            }
        }
        // When segment hasn't been made before
        if (isFound == 0) {
            // Allocate new page table slot
            seg_table->table[seg_table->size].v_index = first_lv;
            seg_table->table[seg_table->size].pages =
                (struct page_table_t*)malloc(sizeof(struct page_table_t));

            // Allocate new page translator in corresponding table
            seg_table->table[seg_table->size].pages->table[0].v_index = second_lv;
            seg_table->table[seg_table->size].pages->table[0].p_index = i;

            seg_table->table[seg_table->size].pages->size = 1;

            seg_table->size++;
        }
        count_alloc_pages++;
        if (count_alloc_pages == num_pages) {
            _mem_stat[i].next = -1;
            break;
        }
    }
}
pthread_mutex_unlock(&mem_lock);

```



```

    return ret_mem;
}

```

2.2.4 Free memory

From transformed address, we find all allocated pages and delete them, then updating page table, if page table is deleted, we update seg table again. The content of `free_mem()`:

```

int free_mem(addr_t address, struct pcb_t* proc) {
    /*TODO: Release memory region allocated by [proc]. The first byte of
     * this region is indicated by [address]. Task to do:
     * - Set flag [proc] of physical page use by the memory block
     * back to zero to indicate that it is free.
     * - Remove unused entries in segment table and page tables of
     * the process [proc].
     * - Remember to use lock to protect the memory from other
     * processes. */
    pthread_mutex_lock(&mem_lock);
    struct page_table_t* page_table = get_page_table(get_first_lv(address),
        proc->seg_table);

    int valid = 0;
    if (page_table != NULL) {
        int i;
        for (i = 0; i < page_table->size; i++) {
            if (page_table->table[i].v_index == get_second_lv(address)) {
                addr_t physical_addr;
                if (translate(address, &physical_addr, proc)) {
                    int p_index = physical_addr >> OFFSET_LEN; // 10 bits page index
                    int num_free_pages = 0;
                    addr_t cur_vir_addr = (num_free_pages << OFFSET_LEN) + address;
                    addr_t seg_idx, page_idx;
                    do {
                        _mem_stat[p_index].proc = 0;
                        int found = 0;
                        int k;
                        seg_idx = get_first_lv(cur_vir_addr);

```

```

page_idx = get_second_lv(cur_vir_addr);
for (k = 0; k < proc->seg_table->size && !found; k++) {
    if (proc->seg_table->table[k].v_index == seg_idx) {
        int l;
        for (l = 0; l < proc->seg_table->table[k].pages
->size; l++) {
            if (proc->seg_table->table[k].pages->table[l].
v_index == page_idx) {
                int m;
                //Rearrange page table
                for (m = l; m < proc->seg_table->table[k].
pages->size - 1; m++)
                    proc->seg_table->table[k].pages->table[m]
                    = proc->seg_table->table[k].pages->table
                    [m+1];
                proc->seg_table->table[k].pages->size--;
                //If page empty
                if(proc->seg_table->table[k].pages->size==0)
                {
                    free(proc->seg_table->table[k].pages);
                    //Rearrange segment table
                    for (m = k; m < proc->seg_table->size - 1;
m++)
                        proc->seg_table->table[m] =
                        proc->seg_table->table[m + 1];
                        proc->seg_table->size--;
                }
                found = 1;
                break;
            }
        }
    }
}

p_index = _mem_stat[p_index].next;
num_free_pages++;
} while (p_index != -1);
valid = 1;

```

```

        }
        break;
    }
}

pthread_mutex_unlock(&mem_lock);
return valid ? 0 : 1;
}

```

2.3 Test run

After compiling `alloc_mem()` and `free_mem()`, we run program by "make mem" trong Terminal. After that, run "make test_mem" to show result:

```

000: 00000-003ff - PID: 01 (idx 000, nxt: 001)
      003e8: 14
001: 00400-007ff - PID: 01 (idx 001, nxt: -01)
002: 00800-00bff - PID: 01 (idx 000, nxt: 003)
003: 00c00-00fff - PID: 01 (idx 001, nxt: 004)
004: 01000-013ff - PID: 01 (idx 002, nxt: 005)
005: 01400-017ff - PID: 01 (idx 003, nxt: 006)
006: 01800-01bff - PID: 01 (idx 004, nxt: -01)
014: 03800-03bff - PID: 01 (idx 000, nxt: 015)
      03814: 64
015: 03c00-03fff - PID: 01 (idx 001, nxt: -01)

```

3 Put all together

To ensure processes We will insert lock mechanism in `read_mem()` and `Write_mem()`:

```

int read_mem(addr_t address, struct pcb_t* proc, BYTE* data) {
    addr_t physical_addr;
    if (translate(address, &physical_addr, proc)) {
        pthread_mutex_lock(&mem_lock);
        *data = _ram[physical_addr];
        pthread_mutex_unlock(&mem_lock);
        return 0;
    }
}

```

/

```
    } else {  
        return 1;  
    }  
}  
}  
int write_mem(addr_t address, struct pcb_t* proc, BYTE data) {  
    addr_t physical_addr;  
    if (translate(address, &physical_addr, proc)) {  
        pthread_mutex_lock(&mem_lock);  
        _ram[physical_addr] = data;  
        pthread_mutex_unlock(&mem_lock);  
        return 0;  
    } else {  
        return 1;  
    }  
}
```

Next, run "make all" to check work by first compiling the whole source code. Result of combination of scheduling and memory management:

```
Time slot    0  
Loaded a process at input/proc/p0, PID: 1  
Time slot    1  
CPU 0: Dispatched process 1  
Loaded a process at input/proc/p1, PID: 2  
Time slot    2  
Time slot    3  
CPU 1: Dispatched process 2  
Loaded a process at input/proc/p1, PID: 3  
Time slot    4  
Loaded a process at input/proc/p1, PID: 4  
Time slot    5  
Time slot    6  
Time slot    7  
CPU 0: Put process 1 to run queue  
CPU 0: Dispatched process 4  
Time slot    8  
Time slot    9  
CPU 1: Put process 2 to run queue
```

/

CPU 1: Dispatched process 3
Time slot 10
Time slot 11
Time slot 12
Time slot 13
CPU 0: Put process 4 to run queue
CPU 0: Dispatched process 1
Time slot 14
Time slot 15
CPU 1: Put process 3 to run queue
CPU 1: Dispatched process 2
Time slot 16
Time slot 17
CPU 0: Processed 1 has finished
CPU 0: Dispatched process 4
Time slot 18
Time slot 19
CPU 1: Processed 2 has finished
CPU 1: Dispatched process 3
Time slot 20
Time slot 21
CPU 0: Processed 4 has finished
CPU 0 stopped
Time slot 22
Time slot 23
CPU 1: Processed 3 has finished
CPU 1 stopped

MEMORY CONTENT:

000: 00000-003ff - PID: 04 (idx 000, nxt: 001)
001: 00400-007ff - PID: 04 (idx 001, nxt: 002)
002: 00800-00bff - PID: 04 (idx 002, nxt: 003)
003: 00c00-00fff - PID: 04 (idx 003, nxt: -01)
004: 01000-013ff - PID: 03 (idx 000, nxt: 005)
005: 01400-017ff - PID: 03 (idx 001, nxt: 006)
006: 01800-01bff - PID: 03 (idx 002, nxt: 012)
007: 01c00-01fff - PID: 02 (idx 000, nxt: 008)
008: 02000-023ff - PID: 02 (idx 001, nxt: 009)

009: 02400-027ff - PID: 02 (idx 002, nxt: 010)
025e7: 0a
010: 02800-02bff - PID: 02 (idx 003, nxt: 011)
011: 02c00-02fff - PID: 02 (idx 004, nxt: -01)
012: 03000-033ff - PID: 03 (idx 003, nxt: -01)
014: 03800-03bff - PID: 04 (idx 000, nxt: 015)
015: 03c00-03fff - PID: 04 (idx 001, nxt: 016)
016: 04000-043ff - PID: 04 (idx 002, nxt: 017)
041e7: 0a
017: 04400-047ff - PID: 04 (idx 003, nxt: 018)
018: 04800-04bff - PID: 04 (idx 004, nxt: -01)
023: 05c00-05fff - PID: 02 (idx 000, nxt: 024)
024: 06000-063ff - PID: 02 (idx 001, nxt: 025)
025: 06400-067ff - PID: 02 (idx 002, nxt: 026)
026: 06800-06bff - PID: 02 (idx 003, nxt: -01)
047: 0bc00-0bfff - PID: 01 (idx 000, nxt: -01)
0bc14: 64
057: 0e400-0e7ff - PID: 03 (idx 000, nxt: 058)
058: 0e800-0ebff - PID: 03 (idx 001, nxt: 059)
059: 0ec00-0efff - PID: 03 (idx 002, nxt: 060)
0ede7: 0a
060: 0f000-0f3ff - PID: 03 (idx 003, nxt: 061)
061: 0f400-0f7ff - PID: 03 (idx 004, nxt: -01)

References

- [1] Operating System Concept - Abraham Silberschatz, Greg Gagne, Peter B. Galvin
- [2] GNU Operating System "<https://gnu.org/>".
- [3] Slides OS Concept - chapter 3, 4, 5, 6, 9, 10